

CS 411 Project Track 1 Stage 3

Group Number: 096

GroupName: OpenSource

Each of our 5 entity tables have at least 1000 rows:

```
mysql> show tables;
+-----+
| Tables_in_opensource |
+-----+
| author                |
| music                 |
| topic                 |
| topic_work            |
| university            |
| work                  |
| work_author           |
+-----+
7 rows in set (0.01 sec)
```

```
mysql> select count(*) from author;
+-----+
| count(*) |
+-----+
|    36996 |
+-----+
1 row in set (0.46 sec)
```

```
mysql> select count(*) from topic;
+-----+
| count(*) |
+-----+
|    2892 |
+-----+
1 row in set (0.24 sec)

mysql> select count(*) from university;
+-----+
| count(*) |
+-----+
|   21607 |
+-----+
1 row in set (0.22 sec)
```

```
mysql> select count(*) from work;
+-----+
| count(*) |
+-----+
|    9969 |
+-----+
1 row in set (0.03 sec)
```

```
mysql> select count(*) from music;
+-----+
| count(*) |
+-----+
|    1000 |
+-----+
1 row in set (0.05 sec)
```

```
mysql> select count(*) from topic_work;
+-----+
| count(*) |
+-----+
|   24830 |
+-----+
1 row in set (0.49 sec)
```

```
mysql> select count(*) from work author;
+-----+
| count(*) |
+-----+
|    9969 |
+-----+
1 row in set (0.01 sec)
```

For these queries, the hard coded values would be input from the user (or parsed data from the user).

Query 1: Shows how many works are released each year for medicine.

```
SELECT work.publication_year,  
       Count(work.work_id)  
FROM   work  
       JOIN topic_work  
         ON topic_work.work_id = work.work_id  
       JOIN topic  
         ON topic_work.topic_id = topic.topic_id  
WHERE  topic.category = 'Medicine'  
GROUP BY work.publication_year  
ORDER BY work.publication_year ASC;
```

	publication_year	COUNT(work.work_id)
▶	1889	1
	1907	1
	1909	1
	1938	1
	1949	5
	1950	1
	1952	1
	1954	1
	1955	1
	1956	5
	1957	4
	1958	2
	1959	6
	1960	1
	1961	4

Query 2: Shows works published by universities other than the one selected

```
SELECT DISTINCT work.work_id,  
                author.university_name  
FROM    work  
        JOIN work_author  
            ON work.work_id = work_author.work_id  
        JOIN author  
            ON author.author_id = work_author.author_id  
        JOIN university  
            ON university.institution_name = author.university_name  
WHERE    university.country = 'United States'  
EXCEPT  
SELECT DISTINCT work.work_id,  
                author.university_name  
FROM    work  
        JOIN work_author  
            ON work.work_id = work_author.work_id  
        JOIN author  
            ON author.author_id = work_author.author_id  
WHERE    author.university_name = 'Cornell University';
```

	work_id	university_name
▶	https://openalex.org/W1518101824	Harvard University
	https://openalex.org/W1513425511	Harvard University
	https://openalex.org/W1487557448	Harvard University
	https://openalex.org/W1004584821	Harvard University
	https://openalex.org/W1514621373	Harvard University Press
	https://openalex.org/W1509292812	Harvard University Press
	https://openalex.org/W1255676896	Harvard University Press
	https://openalex.org/W1533942137	Broad Institute
	https://openalex.org/W1255676896	Broad Institute
	https://openalex.org/W1533942137	Massachusetts General Hospital
	https://openalex.org/W1532325895	Stanford University
	https://openalex.org/W1527927437	Stanford University
	https://openalex.org/W1509562192	Stanford University
	https://openalex.org/W1504303645	Johns Hopkins University
	https://openalex.org/W1494198834	Johns Hopkins University

Query 3: This will get all works related to the given paper.

```
SELECT work_outer.work_id
FROM   work work_outer
JOIN   topic topic_outer
JOIN   topic_work tw_outer
ON     work_outer.work_id = tw_outer.work_id
AND    topic_outer.topic_id = tw_outer.topic_id
WHERE  work_outer.work_id <> 'https://openalex.org/works/W157808733'
AND    topic_outer.category IN
(
    SELECT topic.category
    FROM   work
    JOIN   topic_work
    ON     topic_work.work_id = work.work_id
    JOIN   topic
    ON     topic_work.topic_id = topic.topic_id
    WHERE  work.work_id LIKE '%W950821216%' );
```

	work_id
▶	https://openalex.org/works/W1501145371
	https://openalex.org/works/W1550061415
	https://openalex.org/works/W1556801875
	https://openalex.org/works/W1650728208
	https://openalex.org/works/W1963989017
	https://openalex.org/works/W1965092590
	https://openalex.org/works/W1965752441
	https://openalex.org/works/W1966182479
	https://openalex.org/works/W1966750682
	https://openalex.org/works/W1970127494
	https://openalex.org/works/W1971275758
	https://openalex.org/works/W1971472669
	https://openalex.org/works/W1975708211
	https://openalex.org/works/W1976106600
	https://openalex.org/works/W1979544533

Query 4: This returns how many publications each university has made

```
SELECT    Count(work.work_id),  
          author.university_name  
FROM      work  
JOIN      author  
JOIN      work_author  
ON        work_author.work_id = work.work_id  
AND       work_author.author_id = author.author_id  
WHERE     author.university_name <> 'Unknown'  
GROUP BY  author.university_name  
ORDER BY  Count(work.work_id) DESC;
```

	count(work.work_id)	university_name
►	494	Stanford University
	468	Harvard University
	411	Broad Institute
	384	Massachusetts Institute of Technology
	342	University of Oxford
	339	Washington University in St. Louis
	338	University of Washington
	326	Institute for Health Metrics and Evaluation
	313	Wellcome Sanger Institute
	309	University of California, Berkeley
	282	National Institutes of Health
	272	Google (United States)
	252	University of Michigan–Ann Arbor
	250	University of Cambridge
	221	Johns Hopkins University

DDL Commands

```
create table author (  
    author_id varchar(255),  
    author_name varchar(255),  
    university_name varchar(255),  
    PRIMARY KEY (author_id),  
    FOREIGN KEY (university_name)  
        references university(InstitutionId));
```

```
create table university (  
    InstitutionId int,  
    Institution_Name varchar(255),  
    City varchar(255),  
    Country varchar(255),  
    PRIMARY KEY (InstitutionId));
```

```
create table work (  
    work_id varchar(255),  
    title varchar(1000),  
    publication_year int,  
    primary key (work_id));
```

```
create table topic (  
    topic_id varchar(255),  
    topic_name varchar(255),  
    category varchar(255),  
    primary key (topic_id));
```

```
create table music (  
    MusicID int,  
    filename varchar(255),  
    genre varchar(255),  
    Primary key (MusicID));
```

```
create table topic_work (  
    topic_id varchar(255),  
    work_id varchar(255),  
    PRIMARY KEY (topic_id, work_id),  
    FOREIGN KEY (topic_id) REFERENCES topic(topic_id)  
    FOREIGN KEY (work_id) REFERENCES work(work_id));
```

```
create table topic_music (
    category varchar(255),
    genre varchar(255),
    PRIMARY KEY (category),
    FOREIGN KEY (category) REFERENCES topic(category)
    FOREIGN KEY (genre) REFERENCES music(genre));
```

[illegible]

```

-----+
| -> Sort: 'work' publication_year (actual time=23.332..23.335 rows=78 loops=1)
| -> Table scan on 'temporary' (actual time=23.282..23.290 rows=78 loops=1)
| -> Aggregate using temporary table (actual time=23.280..23.280 rows=78 loops=1)
| -> Nested loop inner join (cost=1425.03 rows=1948) (actual time=0.093..22.209 rows=3354 loops=1)
| -> Nested loop inner join (cost=743.19 rows=1948) (actual time=0.081..10.484 rows=3354 loops=1)
| -> Filter: (topic_category = 'Medicine') (cost=294.95 rows=271) (actual time=0.045..1.495 rows=481 loops=1)
| -> Table scan on topic (cost=294.95 rows=2707) (actual time=0.036..1.096 rows=2892 loops=1)
| -> Covering index lookup on topic_work using PRIMARY (topic_id=topic.topic_id) (cost=0.94 rows=7) (actual time=0.014..0.018 rows=7 loops=
481)
| -> Single-row index lookup on work using PRIMARY (work_id=topic_work.work_id) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=3354)

```

```

| -> Sort: 'work'.publication_year (actual time=19.909..19.914 rows=78 loops=1)
| -> Table scan on <temporary> (actual time=19.867..19.876 rows=78 loops=1)
|   -> Aggregate using temporary table (actual time=19.864..19.864 rows=78 loops=1)
|     -> Nested loop inner join (cost=2145.84 rows=3462) (actual time=0.145..18.861 rows=3354 loops=1)
|       -> Nested loop inner join (cost=234.29 rows=3462) (actual time=0.132..7.955 rows=3354 loops=1)
|         -> Covering index lookup on topic using index2 (category='Medicine') (cost=137.83 rows=481) (actual time=0.107..0.331 rows=481 loops=1)
|         -> Covering index lookup on topic_work using PRIMARY (topic_id=topic.topic_id) (cost=0.94 rows=7) (actual time=0.012..0.015 rows=7 loops=
481)
|       -> Single-row index lookup on work using PRIMARY (work_id=topic_work.work_id) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=3354)
|

```

With index on topic.category

```
| -> Sort: 'work'.publication_year (actual time=21.315..21.319 rows=78 loops=1)
-> Table scan on <temporary> (actual time=21.267..21.276 rows=78 loops=1)
-> Aggregate using temporary table (actual time=21.265..21.265 rows=78 loops=1)
-> Nested loop inner join (cost=2145.84 rows=3462) (actual time=0.076..20.042 rows=3354 loops=1)
-> Nested loop inner join (cost=934.29 rows=3462) (actual time=0.063..7.806 rows=3354 loops=1)
-> Covering index lookup on topic using index2 (category='Medicine') (cost=137.83 rows=481) (actual time=0.035..0.300 rows=481 loops=1)
-> Covering index lookup on topic_work using PRIMARY (topic_id=topic.topic_id) (cost=0.94 rows=7) (actual time=0.011..0.015 rows=7 loops=1)
481)
-> Single-row index lookup on work using PRIMARY (work_id=topic_work.work_id) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=3354)
|
```

We will only keep work.publication_year indexing.

Work.publication_year was kept and very helpful. It is because there are many years in the table, and the year shows up in intensive clauses like GROUP BY and ORDER BY. Both of these take up resources and seem to requery the fields, making indexing worthwhile.

topics.category seems to help as it is used in some intensive SQL clauses, particularly the where clause that checks if it is equal to a given value.

Query 2 without any indexing:

```
-----+
| -> Table scan on <except temporary> (cost=32128.02..32349.97 rows=17557) (actual time=2819.528..2820.778 rows=9085 loops=1)
-> Except materialize with deduplication (cost=32128.01..32128.01 rows=17557) (actual time=2819.526..2819.526 rows=9183 loops=1)
-> Table scan on <temporary> (cost=29947.68..30169.63 rows=17557) (actual time=2798.573..2800.263 rows=9183 loops=1)
-> Temporary table with deduplication (cost=29947.67..29947.67 rows=17557) (actual time=2798.568..2798.568 rows=9183 loops=1)
-> Nested loop inner join (cost=28191.93 rows=17557) (actual time=135.920..2713.103 rows=72714 loops=1)
-> Nested loop inner join (cost=22046.85 rows=17557) (actual time=135.889..2442.864 rows=72714 loops=1)
-> Nested loop inner join (cost=8216.51 rows=15153) (actual time=62.095..842.304 rows=53481 loops=1)
-> Filter: (university.country = 'United States') and (university.institution_name is not null) (cost=2110.29 rows=1951) (actual time=15.099..369.694 rows=5549 loops=1)
-> Table scan on university (cost=2110.29 rows=19505) (actual time=13.881..362.640 rows=21607 loops=1)
-> Covering index lookup on author using idx1 (university_name=university.institution_name) (cost=2.35 rows=8) (actual time=0.071..0.084 rows=10 loops=5549)
-> Covering index lookup on work_author using author_id (author_id=author.author_id) (cost=0.80 rows=1) (actual time=0.028..0.030 rows=1 loops=53481)
-> Single-row covering index lookup on work using PRIMARY (work_id=work_author.work_id) (cost=0.25 rows=1) (actual time=0.003..0.004 rows=1 loops=72714)
-> Table scan on <temporary> (cost=198.52..202.64 rows=132) (actual time=2.629..2.646 rows=98 loops=1)
-> Temporary table with deduplication (cost=198.49..198.49 rows=132) (actual time=2.625..2.625 rows=98 loops=1)
-> Nested loop inner join (cost=185.28 rows=132) (actual time=0.636..2.223 rows=153 loops=1)
-> Nested loop inner join (cost=139.05 rows=132) (actual time=0.619..1.510 rows=153 loops=1)
-> Covering index lookup on author using idx1 (university_name='Cornell University') (cost=35.00 rows=114) (actual time=0.593..0.656 rows=114 loops=1)
-> Covering index lookup on work_author using author_id (author_id=author.author_id) (cost=0.80 rows=1) (actual time=0.006..0.007 rows=1 loops=114)
-> Single-row covering index lookup on work using PRIMARY (work_id=work_author.work_id) (cost=0.25 rows=1) (actual time=0.004..0.004 rows=1 loops=153)
|
-----+-----
```

With index on university.country

```
-----+
| -> Table scan on <except temporary> (cost=92337.35..92964.20 rows=49949) (actual time=1202.613..1203.825 rows=9085 loops=1)
-> Except materialize with deduplication (cost=92337.34..92337.34 rows=49949) (actual time=1202.612..1202.612 rows=9183 loops=1)
-> Table scan on <temporary> (cost=86495.41..87122.26 rows=49949) (actual time=1188.854..1190.486 rows=9183 loops=1)
-> Temporary table with deduplication (cost=86495.40..86495.40 rows=49949) (actual time=1188.850..1188.850 rows=9183 loops=1)
-> Nested loop inner join (cost=81500.48 rows=49949) (actual time=20.162..1113.603 rows=72714 loops=1)
-> Nested loop inner join (cost=34389.33 rows=49949) (actual time=0.896..579.814 rows=72714 loops=1)
-> Nested loop inner join (cost=18047.29 rows=43108) (actual time=0.881..246.538 rows=53481 loops=1)
-> Filter: (university.institution_name is not null) (cost=675.65 rows=5549) (actual time=0.851..19.163 rows=5549 loops=1)
-> Index lookup on university using index3 (Country='United States') (cost=675.65 rows=5549) (actual time=0.848..18.343 rows=5549 loops=1)
-> Covering index lookup on author using idx1 (university_name=university.institution_name) (cost=2.35 rows=8) (actual time=0.033..0.040 rows=10 loops=5549)
-> Covering index lookup on work_author using author_id (author_id=author.author_id) (cost=0.26 rows=1) (actual time=0.005..0.006 rows=1 loops=53481)
-> Single-row covering index lookup on work using PRIMARY (work_id=work_author.work_id) (cost=0.84 rows=1) (actual time=0.007..0.007 rows=1 loops=72714)
-> Table scan on <temporary> (cost=216.04..220.16 rows=132) (actual time=3.194..3.210 rows=98 loops=1)
-> Temporary table with deduplication (cost=216.01..216.01 rows=132) (actual time=3.190..3.190 rows=98 loops=1)
-> Nested loop inner join (cost=202.80 rows=132) (actual time=1.521..3.026 rows=153 loops=1)
-> Nested loop inner join (cost=78.22 rows=132) (actual time=1.508..2.425 rows=153 loops=1)
-> Covering index lookup on author using idx1 (university_name='Cornell University') (cost=35.00 rows=114) (actual time=1.475..1.539 rows=114 loops=1)
-> Covering index lookup on work_author using author_id (author_id=author.author_id) (cost=0.26 rows=1) (actual time=0.006..0.008 rows=1 loops=114)
-> Single-row covering index lookup on work using PRIMARY (work_id=work_author.work_id) (cost=0.84 rows=1) (actual time=0.004..0.004 rows=1 loops=153)
|
-----+-----
```

With index on author.university_name


```

| -> Table scan on <except temporary> (cost=41178.93..41400.88 rows=17557) (actual time=2531.956..2533.415 rows=9085 loops=1)
-> Except materialize with deduplication (cost=41178.92..41178.92 rows=17557) (actual time=2531.955..2531.955 rows=9183 loops=1)
-> Table scan on <temporary> (cost=38930.11..39152.06 rows=17557) (actual time=2517.016..2518.944 rows=9183 loops=1)
-> Temporary table with deduplication (cost=38930.10..38930.10 rows=17557) (actual time=2517.013..2517.013 rows=9183 loops=1)
-> Nested loop inner join (cost=25174.36 rows=17557) (actual time=46.928..2431.496 rows=72714 loops=1)
-> Nested loop inner join (cost=23846.72 rows=17557) (actual time=46.897..1885.567 rows=72714 loops=1)
-> Nested loop inner join (cost=8096.97 rows=15153) (actual time=43.868..256.506 rows=53481 loops=1)
-> Filter: (university.country = 'United States') and (university.institution_name is not null) (cost=1990.75 rows=1951) (actual time=33.263..51.053 rows=5549 loops=1)
-> Table scan on university (cost=1990.75 rows=19505) (actual time=30.750..44.772 rows=21607 loops=1)
-> Covering index lookup on author using idx1 (university_name=university.institution_name) (cost=2.35 rows=8) (actual time=0.027..0.036 rows=10 loops=5549)
-> Covering index lookup on work_author using author_id (author_id=author.author_id) (cost=0.92 rows=1) (actual time=0.029..0.030 rows=1 loops=53481)
-> Single-row covering index lookup on work using PRIMARY (work_id=work_author.work_id) (cost=0.66 rows=1) (actual time=0.007..0.007 rows=1 loops=72714)
-> Table scan on <temporary> (cost=267.00..271.12 rows=132) (actual time=2.134..2.158 rows=98 loops=1)
-> Temporary table with deduplication (cost=266.97..266.97 rows=132) (actual time=2.131..2.131 rows=98 loops=1)
-> Nested loop inner join (cost=251.76 rows=132) (actual time=0.115..1.194 rows=153 loops=1)
-> Nested loop inner join (cost=153.49 rows=132) (actual time=0.101..1.213 rows=153 loops=1)
-> Covering index lookup on author using idx1 (university_name='Cornell University') (cost=35.00 rows=114) (actual time=0.078..0.160 rows=114 loops=1)
-> Covering index lookup on work_author using author_id (author_id=author.author_id) (cost=0.92 rows=1) (actual time=0.008..0.009 rows=1 loops=114)
-> Single-row covering index lookup on work using PRIMARY (work_id=work_author.work_id) (cost=0.66 rows=1) (actual time=0.005..0.005 rows=1 loops=153)
|

```

We did not keep university.country because it made performance worse. It may be because there are very few countries relative to the number of entries in the university table. The overhead of maintaining indexing is not worth it for such few performance gains.

This didn't get us much. This may be because many author.university_name entries are equal to Unknown so we don't get as much benefit. This field also isn't used in any intensive SQL clause so its overhead isn't worth it.

Query 3 without any indexing:

```

| -> Nested loop inner join (cost=8500966.69 rows=25901243) (actual time=317.009..631.355 rows=1315 loops=1)
-> Nested loop inner join (cost=5908407.05 rows=25901243) (actual time=316.973..622.445 rows=1315 loops=1)
-> Nested loop inner join (cost=720389.49 rows=7197533) (actual time=313.781..317.915 rows=93 loops=1)
-> Filter: (topic_outer.category is not null) (cost=365.50 rows=2707) (actual time=47.930..144.985 rows=2892 loops=1)
-> Covering index scan on topic_outer using index2 (cost=365.50 rows=2707) (actual time=47.927..144.617 rows=2892 loops=1)
-> Single-row index lookup on <subquery2> using <auto_distinct_key> (category=topic_outer.category) (actual time=0.060..0.060 rows=0 loops=2892)
-> Materialize with deduplication (cost=6014.41..6014.41 rows=2659) (actual time=170.849..170.849 rows=1 loops=1)
-> Filter: (topic.category is not null) (cost=5748.53 rows=2659) (actual time=170.804..170.828 rows=1 loops=1)
-> Nested loop inner join (cost=5748.53 rows=2659) (actual time=170.801..170.825 rows=1 loops=1)
-> Nested loop inner join (cost=2884.21 rows=2659) (actual time=131.251..131.274 rows=1 loops=1)
-> Filter: ('work'.work_id like '%W950821216%') (cost=1016.45 rows=1085) (actual time=131.205..131.219 rows=1 loops=1)
-> Covering index scan on work using PRIMARY (cost=1016.45 rows=9762) (actual time=0.057..126.320 rows=9969 loops=1)
-> Covering index lookup on topic_work using work_id (work_id='work'.work_id) (cost=1.48 rows=2) (actual time=0.044..0.052 rows=1 loops=1)
-> Single-row index lookup on topic using PRIMARY (topic_id=topic_work.topic_id) (cost=0.98 rows=1) (actual time=39.548..39.548 rows=1 loops=1)
-> Filter: (tw_outer.work_id <> 'https://openalex.org/works/W157808733') (cost=3.05 rows=4) (actual time=2.599..3.273 rows=14 loops=93)
-> Covering index lookup on tw_outer using PRIMARY (topic_id=topic_outer.topic_id) (cost=3.05 rows=7) (actual time=2.596..3.262 rows=14 loops=93)
-> Single-row covering index lookup on work_outer using PRIMARY (work_id=tw_outer.work_id) (cost=0.25 rows=1) (actual time=0.007..0.007 rows=1 loops=1315)
|

```

With index on topic.category

```

| -> Nested loop inner join (cost=8498243.97 rows=25901243) (actual time=9.744..17.600 rows=1315 loops=1)
-> Nested loop inner join (cost=5905684.33 rows=25901243) (actual time=9.720..13.017 rows=1315 loops=1)
-> Nested loop inner join (cost=720318.94 rows=7197533) (actual time=9.676..10.176 rows=93 loops=1)
-> Filter: (topic_outer.category is not null) (cost=294.95 rows=2707) (actual time=0.039..0.975 rows=2892 loops=1)
-> Covering index scan on topic_outer using index2 (cost=294.95 rows=2707) (actual time=0.038..0.777 rows=2892 loops=1)
-> Single-row index lookup on <subquery2> using <auto_distinct_key> (category=topic_outer.category) (actual time=0.003..0.003 rows=0 loops=2892)
-> Materialize with deduplication (cost=3998.74..3998.74 rows=2659) (actual time=7.608..7.608 rows=1 loops=1)
-> Filter: (topic.category is not null) (cost=3732.86 rows=2659) (actual time=7.585..7.590 rows=1 loops=1)
-> Nested loop inner join (cost=3732.86 rows=2659) (actual time=7.582..7.587 rows=1 loops=1)
-> Nested loop inner join (cost=2802.25 rows=2659) (actual time=7.559..7.564 rows=1 loops=1)
-> Filter: ('work'.work_id like '%W950821216%') (cost=1016.45 rows=1085) (actual time=7.506..7.508 rows=1 loops=1)
-> Covering index scan on work using PRIMARY (cost=1016.45 rows=9762) (actual time=0.059..2.848 rows=9969 loops=1)
-> Covering index lookup on topic_work using work_id (work_id='work'.work_id) (cost=1.40 rows=2) (actual time=0.050..0.052 rows=1 loops=1)
-> Single-row index lookup on topic using PRIMARY (topic_id=topic_work.topic_id) (cost=0.25 rows=1) (actual time=0.021..0.021 rows=1 loops=1)
-> Filter: (tw_outer.work_id <> 'https://openalex.org/works/W157808733') (cost=2.07 rows=4) (actual time=0.016..0.030 rows=14 loops=93)
-> Covering index lookup on tw_outer using PRIMARY (topic_id=topic_outer.topic_id) (cost=2.07 rows=7) (actual time=0.015..0.023 rows=14 loops=93)
-> Single-row covering index lookup on work_outer using PRIMARY (work_id=tw_outer.work_id) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=1315)
|

```

topic.category is the only field in this query that we should manually apply an index to as it is the only non-primary key. However, it does not seem to have any effect on the cost. This may be because there are relatively few categories compared to topics so there aren't many benefits to indexing this.

Query 4 without any indexing

```
| -> Sort: Count(work.work_id) DESC (actual time=1518.270..1518.626 rows=4020 loops=1)
    -> Stream results (cost=25187.82 rows=22135) (actual time=200.797..1511.871 rows=4020 loops=1)
        -> Group aggregate: count('work'.work_id) (cost=25187.82 rows=22135) (actual time=200.791..1508.643 rows=4020 loops=1)
            -> Nested loop inner join (cost=22974.37 rows=22135) (actual time=88.964..1469.238 rows=35734 loops=1)
                -> Nested loop inner join (cost=15227.28 rows=22135) (actual time=88.931..1254.138 rows=35734 loops=1)
                    -> Filter: (author.university_name <> 'Unknown') (cost=3918.13 rows=19103) (actual time=0.151..220.174 rows=26965 loops=1)
                    -> Covering index scan on author using idx1 (cost=3918.13 rows=35021) (actual time=0.143..208.645 rows=36996 loops=1)
                    -> Covering index lookup on work_author using author_id (author_id=author.author_id) (cost=0.48 rows=1) (actual time=0.036..0.038 rows=1 loops=26965)
                -> Single-row covering index lookup on work using PRIMARY (work_id=work_author.work_id) (cost=0.25 rows=1) (actual time=0.006..0.006 rows=1 loops=35734)
            |
```

With index of author.university_name

```
| -> Sort: Count(work.work_id) DESC (actual time=1247.394..1247.942 rows=4020 loops=1)
    -> Stream results (cost=38054.37 rows=22135) (actual time=46.058..1245.565 rows=4020 loops=1)
        -> Group aggregate: count('work'.work_id) (cost=38054.37 rows=22135) (actual time=46.049..1243.338 rows=4020 loops=1)
            -> Nested loop inner join (cost=35840.92 rows=22135) (actual time=39.216..1223.381 rows=35734 loops=1)
                -> Nested loop inner join (cost=11643.84 rows=22135) (actual time=0.074..655.088 rows=35734 loops=1)
                    -> Filter: (author.university_name <> 'Unknown') (cost=3629.73 rows=19103) (actual time=0.048..164.047 rows=26965 loops=1)
                    -> Covering index scan on author using idx1 (cost=3629.73 rows=35021) (actual time=0.042..156.963 rows=36996 loops=1)
                    -> Covering index lookup on work_author using author_id (author_id=author.author_id) (cost=0.30 rows=1) (actual time=0.017..0.018 rows=1 loops=26965)
                -> Single-row covering index lookup on work using PRIMARY (work_id=work_author.work_id) (cost=0.99 rows=1) (actual time=0.016..0.016 rows=1 loops=35734)
            |
```

This didn't get us much. This may be because many author.university_name entries are equal to Unknown so we don't get as much benefit. This field also isn't used in any intensive SQL clause so its overhead isn't worth it.

We will not be keeping this indexing.

For queries 3 and 4, we only have one field that is not a primary key. Therefore, we can only try one configuration (other than the default).